

## reduce(func)

Aggregate the elements of a dataset through *func*

```
>>> names1 = sc.parallelize(["abe", "abby", "apple"])
>>> print names1.reduce(lambda t1, t2: t1+t2)
abeabbyapple

>>> names2 = sc.parallelize(["apple", "beatty", "beatrice"]).map(lambda a:
[a, len(a)])
>>> print names2.collect()
[['apple', 5], ['beatty', 6], ['beatrice', 8]]

>>> names2.flatMap(lambda t: [t[1]]).reduce(lambda t1, t2: t1+t2)
19
```

## collect(func)

*collect* returns the elements of the RDD back to the driver program.

*collect* is often used in previously provided examples such as [Spark Transformation Examples in Python](#) in order to show the values of the return. Pyspark, for example, will print the values of the array back to the console. This can be helpful in debugging programs.

Examples

```
>>> sc.parallelize([1,2,3]).flatMap(lambda x: [x,x,x]).collect()
[1, 1, 1, 2, 2, 2, 3, 3, 3]
```

## count()

Number of elements in the RDD

```
>>> names1 = sc.parallelize(["abe", "abby", "apple"])
>>> names1.count()
3
```

## first()

Return the first element in the RDD

```
>>> names1 = sc.parallelize(["abe", "abby", "apple"])
>>> names1.first()
'abe'
```

## take(n)

Take the first  $n$  elements of the RDD.

Works by first scanning one partition, and use the results from that partition to estimate the number of additional partitions needed to satisfy the limit.

Translated from the Scala implementation in `RDD#take()`.

Can be much more convenient and economical to use *take* instead of *collect* to inspect a very large RDD

```
>>> names1 = sc.parallelize(["abe", "abby", "apple"])
>>> names1.take(2)
['abe', 'abby']
```

## takeSample(withReplacement, n, seed=None)

Similar to *take*, in return size of  $n$ . Includes boolean option of with or without replacement and random generator seed which defaults to None

```
>>> teams = sc.parallelize(("twins", "brewers", "cubs", "white sox",
"indians", "bad news bears"))
>>> teams.takeSample(True, 3)
['brewers', 'brewers', 'twins']
# run a few times to see different results
```

## countByKey()

Count the number of elements for each key, and return the result to the master as a dictionary.

```
>>> hockeyTeams = sc.parallelize(("wild", "blackhawks", "red wings",
"wild", "oilers", "whalers", "jets", "wild"))
>>> hockeyTeams.map(lambda k: (k,1)).countByKey().items()
[('red wings', 1),
 ('oilers', 1),
 ('blackhawks', 1),
 ('jets', 1),
 ('wild', 3),
 ('whalers', 1)]
```

## saveAsTextFile(path, compressionCodecClass=None)

Save RDD as text file, using string representations of elements.

### Parameters:

- **path** – path to file

- **compressionCodecClass** – (None by default) string i.e. “org.apache.hadoop.io.compress.GzipCodec”

```
>>> hockeyTeams = sc.parallelize(("wild", "blackhawks", "red wings",  
"wild", "oilers", "whalers", "jets", "wild"))  
>>> hockeyTeams.map(lambda k: (k,1)).countByKey().items()  
>>> hockeyTeams.saveAsTextFile("hockey_teams.txt")
```

Produces:

```
$ ls hockey_teams.txt/  
_SUCCESS      part-00001      part-00003      part-00005      part-00007  
part-00000     part-00002      part-00004      part-00006
```

So, you'll see each partition is written to it's own file. I have 8 partitions in dataset example here.